

03 Algorithms and Theory Comparison

Disclosure: all results use the synthetic `fra_synthetic_claims` dataset. Scores may reflect generator rules. Real-world performance is unknown. This demo must not drive real claim decisions.

Summary

Baselines, bagging, boosting (XGBoost, LightGBM, CatBoost), stacking and voting were trained under identical folds. See `benchmark_master.csv` for measured results.

Disclosure: all results use the synthetic `fra_synthetic_claims` dataset. Scores may reflect generator rules. Real-world performance is unknown. This demo must not drive real claim decisions.