

09 Viva Cheat Sheet

Disclosure: all results use the synthetic fra_synthetic_claims dataset. Scores may reflect generator rules. Real-world performance is unknown. This demo must not drive real claim decisions.

Likely questions

Q1. Bias vs variance? A: Single trees high variance; bagging reduces variance; boosting reduces bias but can overfit; we compare families on S1.

Q2. Why grouped splits? A: Claims cluster by district; random CV leaks geography and inflates scores.

Q3. Why not rejection_reason? A: Non-null only when Rejected: perfect target leakage for status.

Q4. Why Set B looks better? A: Gap 0.18053581159110665. Process features are partly consequences of outcome (committee level, processing days).

Q5. Why ship LightGBM if XGBoost wins S1? A: TreeExplainer + Render size. Leaderboard champion is separate from the Flask contract artifact.

Q6. What is calibration? A: Predicted probabilities should match empirical frequencies; we also report abstention and optional conformal sets.

Q7. Why OOF for stacking? A: Meta-learner must not see in-fold base predictions of the same rows.

Q8. Synthetic data? A: Scores may mirror generator rules; must not decide real claims.

Q9. ETA on Pending? A: Trained on Approved/Rejected only; Pending days may be right-censored.

Q10. Fairness? A: Monitoring on synthetic groups only; never automate rejections.